

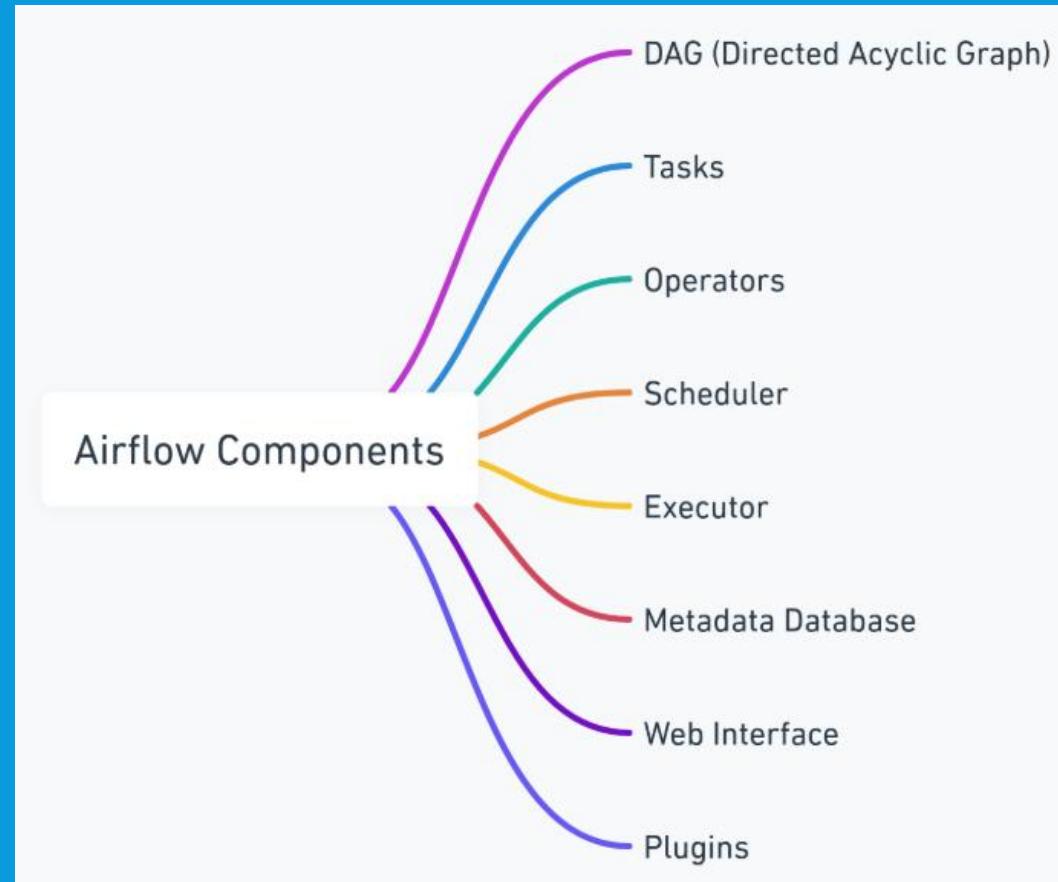


# APACHE-AIRFLOW

# WHAT IS AIRFLOW

- **Apache Airflow is an open-source platform used to programmatically create, schedule, and monitor workflows.**
- **Purpose: It automates complex workflows, making it easy to manage tasks that depend on each other.**
- **Example of a washing machine**

# COMPONENTS OF AIRFLOW




# WHAT IS A WORKFLOW

- A workflow is a **sequence of tasks** or processes that are executed in a specific order to achieve a goal.
- **Steps or Tasks:** Individual components (e.g., downloading data, cleaning it, analyzing it) that make up the workflow.
- **Dependencies:** Tasks are executed based on their relationships or order of execution (e.g., Task B starts only after Task A finishes).
- **Automated Execution:** In Airflow, workflows are automated to save time and reduce manual effort.
- Example - Fetch data from an API → Clean the data → Load it into a database → Generate a report.

# DAG (DIRECTED ACYCLIC GRAPH)

- A DAG is a structure that represents a workflow in Apache Airflow.
- It is a collection of tasks arranged in a specific order
- Tasks don't loop back to previous tasks, meaning no cycles are allowed in the graph.



# OPERATOR

- Each task in a DAG has an Operator that defines what the task needs to do.
- An Operator acts as an instruction for the task.
- It specifies the action the task will perform (e.g., running Python code, executing a SQL query, or waiting for a condition to be met).
- If a task is to run a Python function, you use the PythonOperator.
- PostgresOperator

# SCHEDULER

- The scheduler decides when to start tasks based on their schedule intervals
- Continuously monitors all active DAGs to determine which tasks need to run.
- Ensures tasks are executed in the correct order by checking dependencies defined in the DAG.
- Reschedules tasks that fail, based on retry policies defined in the DAG.

# EXECUTER

- The Executor in Apache Airflow is a component responsible for executing the tasks that the Scheduler schedules. It defines how and where the tasks will run (locally, in parallel, or distributed across multiple systems).
- Executes the tasks that are ready to run based on the scheduler's instructions.
- Manages the allocation of resources (CPU, memory) for task execution.



# METADATA DATABASE

- The **Metadata Database** is a critical component of Apache Airflow. It stores all the information about workflows (DAGs), tasks, and their execution states. Without it, Airflow cannot track or manage workflows effectively.
- Details about DAGs, tasks, schedules, and dependencies.
- The **Scheduler**, **Executor**, and **Web Interface** interact with the metadata database to Read the current state of workflows, Update task statuses, Retrieve historical execution logs for monitoring.

# KEY FEATURES OF APACHE AIRFLOW

- **Scalability:** Can handle workflows of any size.
- **Dynamic :** Workflows are written in Python, allowing dynamic creation and modification.
- **Extensibility:** Provides many plugins and allows you to integrate with external systems like databases, cloud services, and APIs.
- **Monitoring:** Offers a web-based user interface to track progress, logs, and troubleshoot issues.

# TASK LIFE CYCLE

up\_for\_retry upstream\_failed no\_status deferred failed queued removed restarting running scheduled shutdown skipped success up\_for\_reschedule

**up\_for\_retry:** The task failed and is waiting to retry based on the retry policy.

**upstream\_failed:** The task didn't execute because an upstream task failed.

**no\_status:** The task hasn't been assigned any status yet.

**deferred:** The task is waiting for an external trigger to proceed.

**failed:** The task execution failed.

**queued:** The task is waiting in the queue for a worker to pick it up.

**removed:** The task has been removed due to changes in the DAG structure.

**restarting:** The task is restarting after failure or manual intervention.

**running:** The task is currently being executed.

**scheduled:** The task is ready to be executed.

**shutdown:** The task failed and the DAG execution has stopped.

**skipped:** The task was intentionally skipped due to branching or conditions.

**success:** The task was successfully completed.

**up\_for\_reschedule:** The task is waiting for a specific interval to retry execution.

## 1. No Status:

- This is the initial state of a task. It doesn't yet have any defined state in the DAG lifecycle.

## 2. Scheduler:

- The **Scheduler** determines which tasks are ready to run based on their **schedule interval** and **dependencies**.
- The task transitions to **Schedule** if all its upstream tasks are successful or skipped.

## 3. Scheduled:

- The task is now ready to be executed and awaits assignment to an **Executor**.
- From this state, the task may:
  - Proceed to **Executor** to begin execution.
  - Be marked **Removed** if the DAG structure changes.
  - Be marked **Skipped** if a condition (like branching logic) skips the task.
  - Move to **Upstream Failed** if a dependency task fails.

## 4. Executor:

- The **Executor** is responsible for picking up the task and assigning it to a **Worker** for execution.
- The task transitions to **Queued** while it waits for a worker to become available.

## 5. Queued:

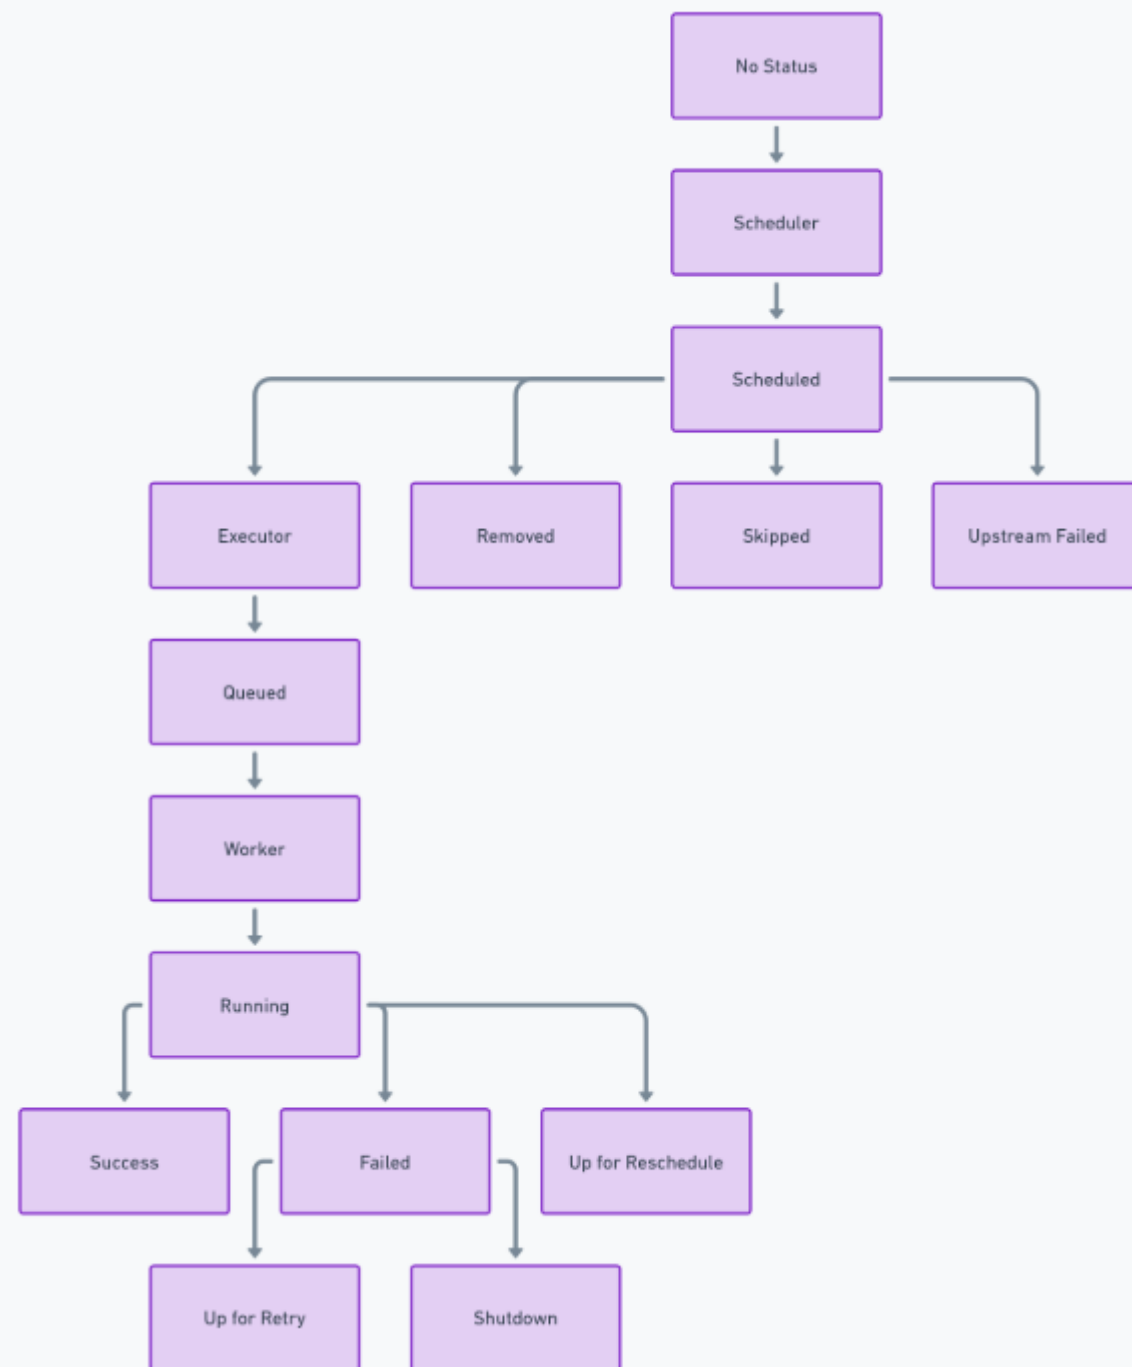
- The task waits in this state until a **Worker** can execute it.
- Once a worker is ready, the task moves to the **Running** state.

## 6. Worker:

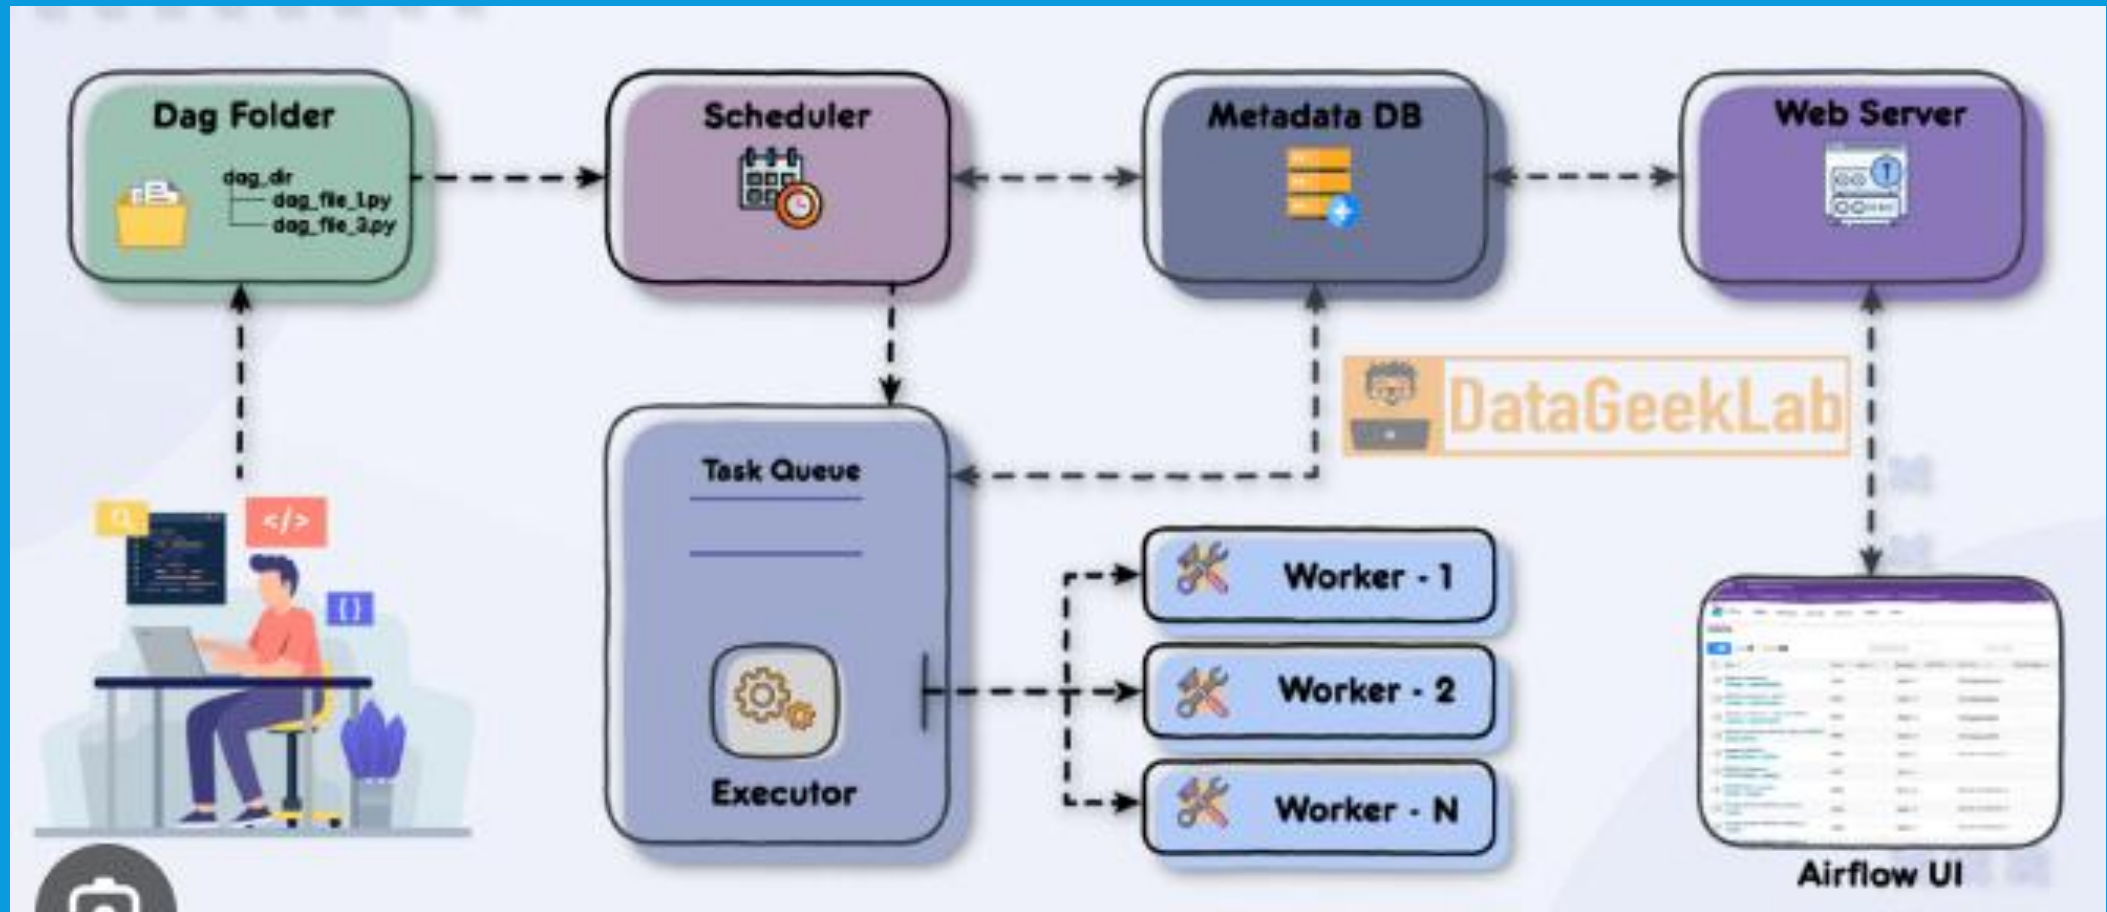
- A **Worker** picks up the task from the queue and begins execution.
- The task status changes to **Running** during this phase.

## 7. Running:

- The task is actively executing.
- After execution:
  - If the task completes successfully, it transitions to **Success**.
  - If the task fails, it transitions to **Failed**.
  - If retries are allowed, it moves to **Up for Retry**.
  - If it is a rescheduled task (e.g., sensors), it moves to **Up for Reschedule**.



# AIRFLOW ARCHITECTURE



# Installing Apache-airflow in windows

# INSTALLING APACHE INSIDE DOCKER

- Install docker and docker compose

```
PS C:\Users\jabcd\OneDrive\Desktop\Apache-Airflow> docker --version
Docker version 27.2.0, build 3ab4256
PS C:\Users\jabcd\OneDrive\Desktop\Apache-Airflow> docker-compose --version
Docker Compose version v2.29.2-desktop.2
PS C:\Users\jabcd\OneDrive\Desktop\Apache-Airflow> |
```

- `curl "https://airflow.apache.org/docs/apache-airflow/2.10.4/docker-compose.yaml" -o 'docker-compose.yaml'`

`mkdir ./dags`

`mkdir ./logs`

`mkdir ./plugins`

`mkdir ./config`

- docker compose up airflow-init

(The docker-compose up airflow-init command initializes Apache Airflow by setting up the **metadata database**, applying required **migrations**, and creating a default **admin user** for the web interface. It ensures the backend is ready for Airflow's scheduler and webserver to function properly. This step is typically required **once** during the initial setup or after resetting the database. You have to run this command only one time The airflow-init service needs to run **once** before starting Airflow for the first time. You only run this command again if, The database is deleted or reset. You make changes to the database configuration. )

- (docker compose up) / (docker compose up -d)
- Without -d, you monitor services interactively but cannot use the same terminal for other commands. With -d, the services run in the background, freeing the terminal for other tasks
- ( docker compose down )/ (docker-compose down -v )

Command	Containers	Networks	Volumes	Use Case
<code>docker-compose down</code>	Removed	Removed	Retained	To preserve data for reuse.
<code>docker-compose down -v</code>	Removed	Removed	Deleted	To reset or start fresh.

Solution: Run Command Prompt as Administrator

1. Run Command Prompt with Elevated Privileges:

- Press `Win + S` and type `cmd`.
- Right-click on Command Prompt and select Run as Administrator.

2. Stop Jenkins:

- Run the following command again:

```
cmd
net stop jenkins
```

Copy code

Airflow username and password is - airflow

- Without -v:** If you stop Airflow using docker-compose down, your **Metadata DB volume** persists, so you can restart later without losing data.
- With -v:** If you use docker-compose down -v, the **Metadata DB volume** is deleted, and you'll need to reinitialize Airflow with docker-compose up airflow-init.



# APACHE-AIRFLOW

- Remove all the Examples DAGS

# REMOVE ALL THE EXAMPLES DAGS

- `docker-compose down -v`

```
environment:
  &airflow-common-env
  AIRFLOW__CORE__EXECUTOR: CeleryExecutor
  AIRFLOW__DATABASE__SQL_ALCHEMY_CONN: postgresql+psycopg2://airflow:airflow@postgres/airflow
  AIRFLOW__CELERY__RESULT_BACKEND: db+postgresql://airflow:airflow@postgres/airflow
  AIRFLOW__CELERY__BROKER_URL: redis://:@redis:6379/0
  AIRFLOW__CORE__FERNET_KEY: ''
  AIRFLOW__CORE__DAGS_ARE_PAUSED_AT_CREATION: 'true'
  AIRFLOW__CORE__LOAD_EXAMPLES: 'true'
  AIRFLOW__API__AUTH_BACKENDS: 'airflow.api.auth.backend.basic_auth,airflow.api.auth.backend.session'
  # yamllint disable rule:line-length
  # Use simple http server as scheduler for health checks
```

`docker-compose up -d`



# APACHE-AIRFLOW

- First DAG from Scratch

# ML END TO END PIPELINE

**1 – EXTRACTING THE DATA**

**2- TRANSFORMING THE DATA**

**3- TRAINING THE MODEL**

**4 – TESTING THE MODEL**



